PREPARED FOR

Emirates

SUBMITTED BY

Aonami

EMIRATES GROUP · RIGHT FIRST TIME (P2P)

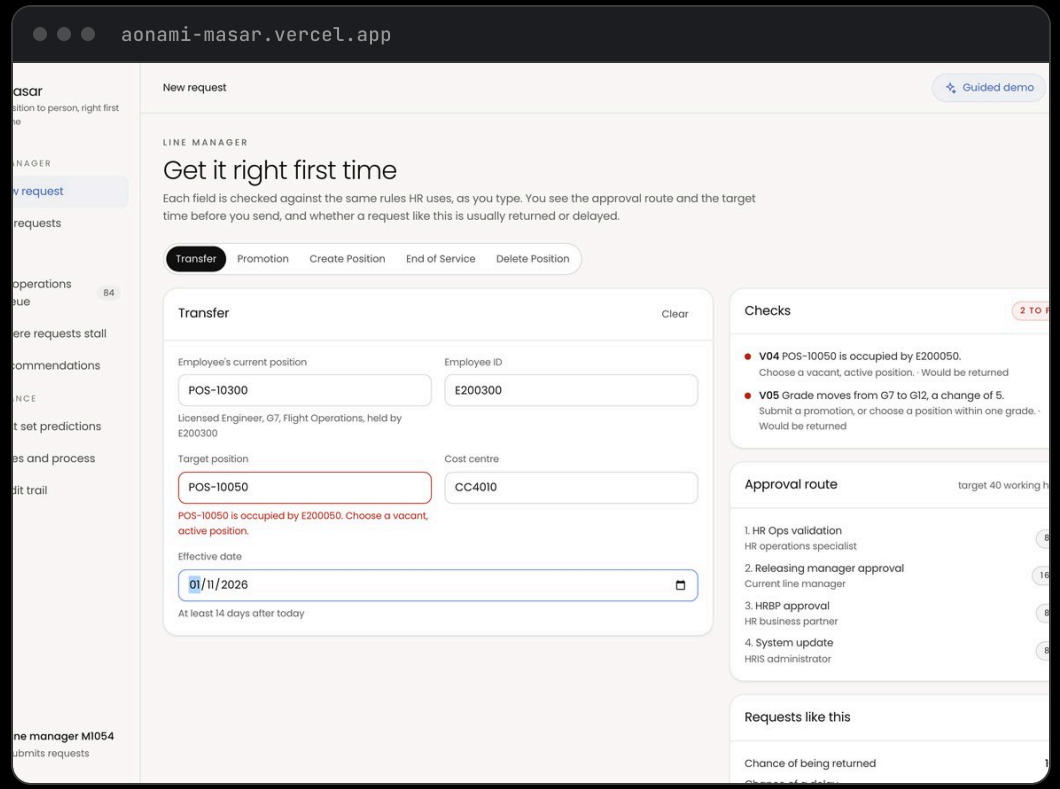
Masar

Position to Person requests, right first time.

LIVE PROTOTYPE

aonami-masar.vercel.app

Press Guided demo. Runs on synthetic sample data.



Requests come back, and nobody can see why.

A Position to Person request passes through HR operations, managers, HR partners, finance and HRIS. Missing or wrong details send it back, dependent transactions hold it, and managers chase by email.

RETURNED AT LEAST ONCE

35%

of requests in our synthetic event log.

LATE

30%

over 125% of the target route time, same log.

VALIDATION RULES

19

checked live in the manager's form.

AT EMIRATES' SCALE

The Emirates Group workforce grew 8% to 130,919 people in 2025-26, with more than 4,000 UAE nationals. [4]

HOW THE CHALLENGE IS JUDGED

JUDGING CRITERION	WEIGHT
Errors caught and prediction accuracy	30%
Process insight	25%
Manager and HR experience	15%
Responsible AI and security	15%
Deployability and reuse	15%

Source: challenge brief [1]

One loop, from submission to insight.

The validation rules become checks in the manager's form; the process map gives the route and target times; the event log shows where requests stall and loop. A transparent predictor scores in-flight requests and lists its evidence. No cloud AI.

1 **Submit**
Each field checked against its rule, with the fix.

2 **Route**
Approvers and target times shown before sending.

3 **Track**
Status, next owner, waiting against target.

4 **Predict**
Return and delay risk per in-flight request, with reasons.

5 **Mine**
Bottlenecks, loops, variants, team variation.

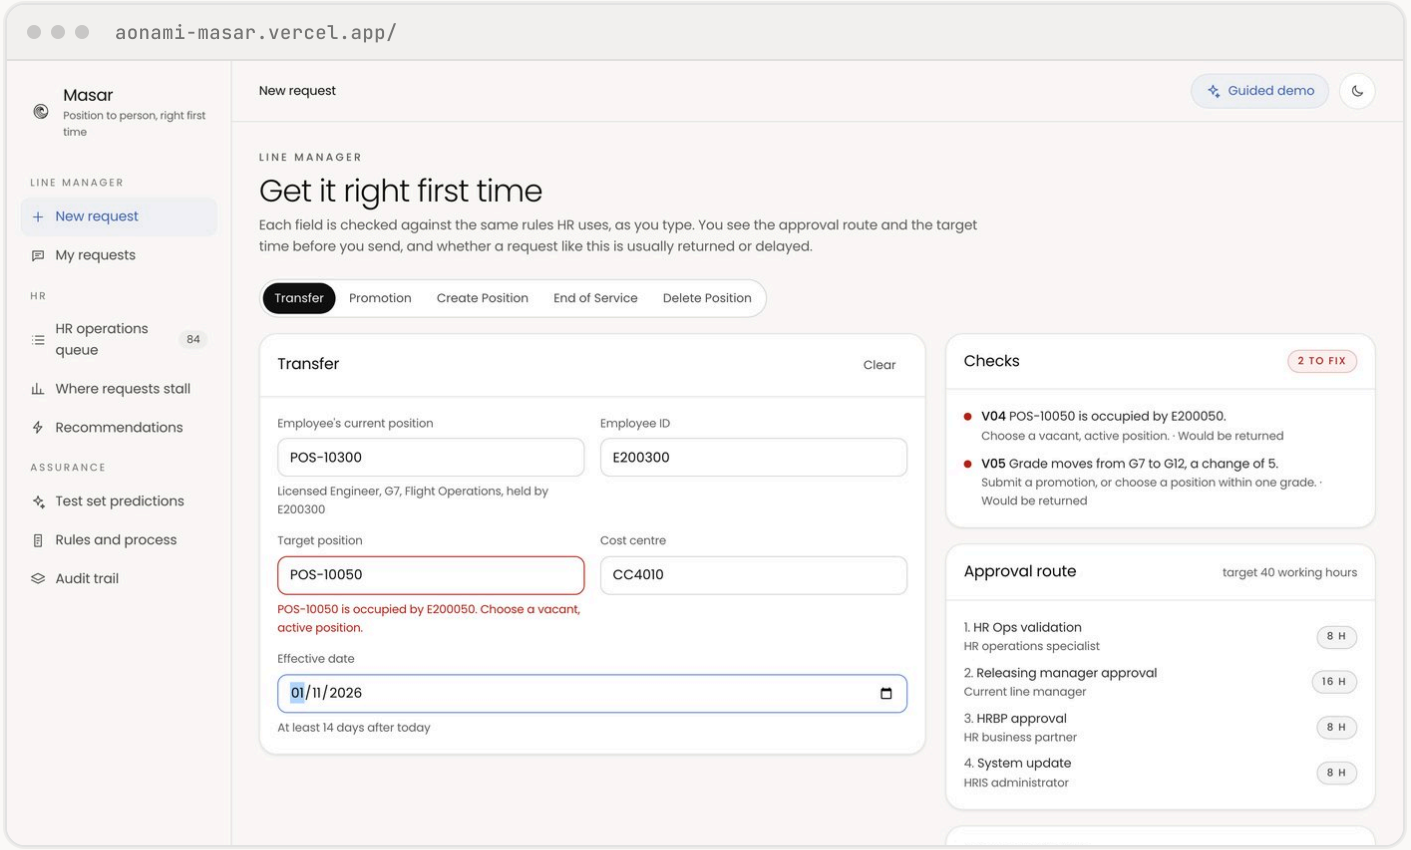
6 **Recommend**
Specific fixes with requests affected and days saved.

- **Masar never approves or returns** a request.
- **Every prediction lists its evidence.**
- **About requests and steps**, never about people.

AI WHERE IT HELPS

Checks and routing are rules; the predictor is a transparent weighting that lists its evidence. Optional: an on-premise model could draft the wording of "where is my request?" replies; it would never approve, return or predict.

Two errors, caught before the request is sent.



V04: target not available

POS-10050 is occupied. Choose a vacant, active position.

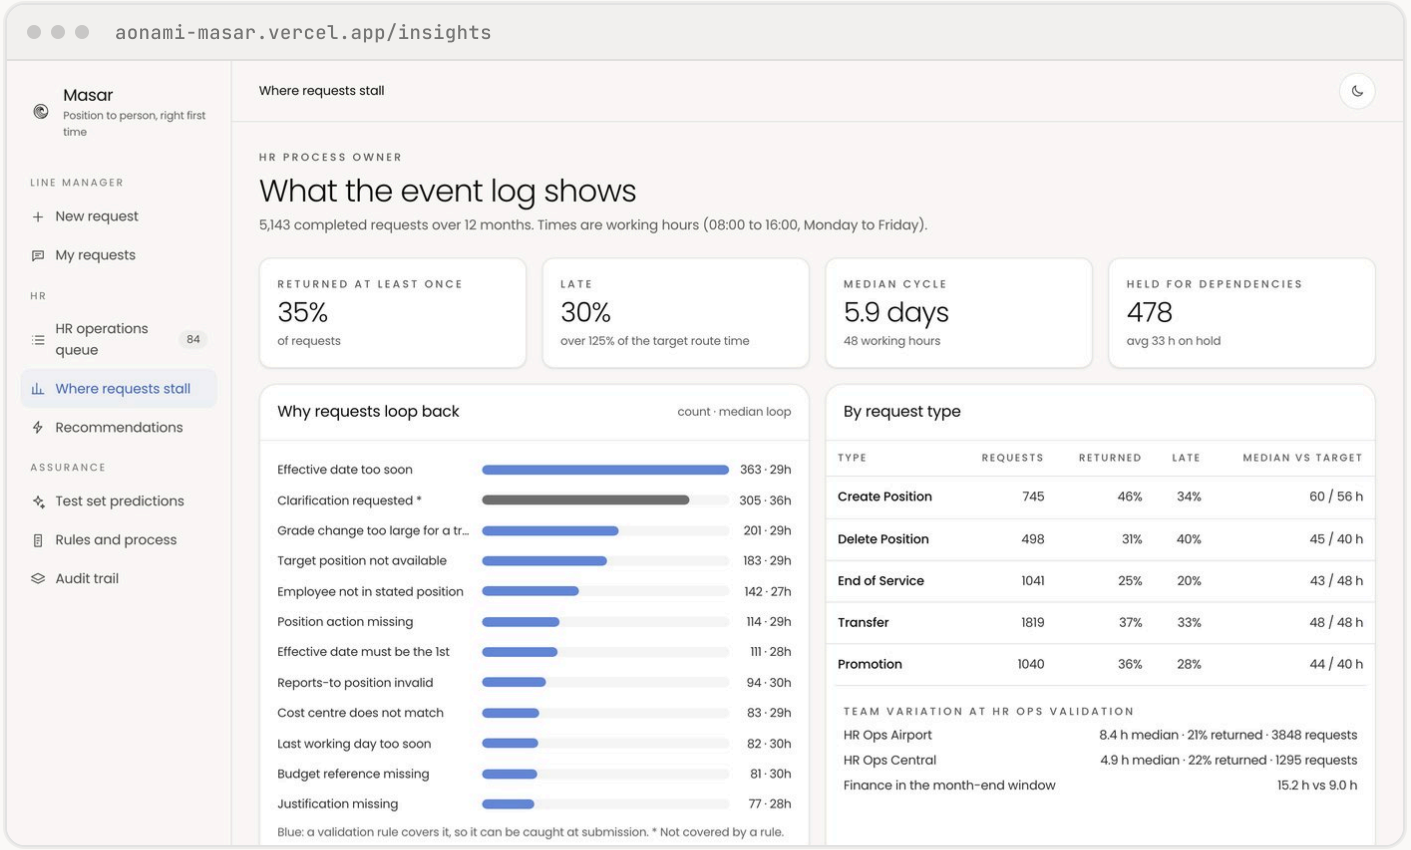
V05: grade change too large

G7 to G12 is a promotion, not a transfer.

Route shown first

Approvers, target time, and how often requests like this come back.

What the event log shows, and what to fix first.



Why requests loop back

Each return reason with its count and median cost.

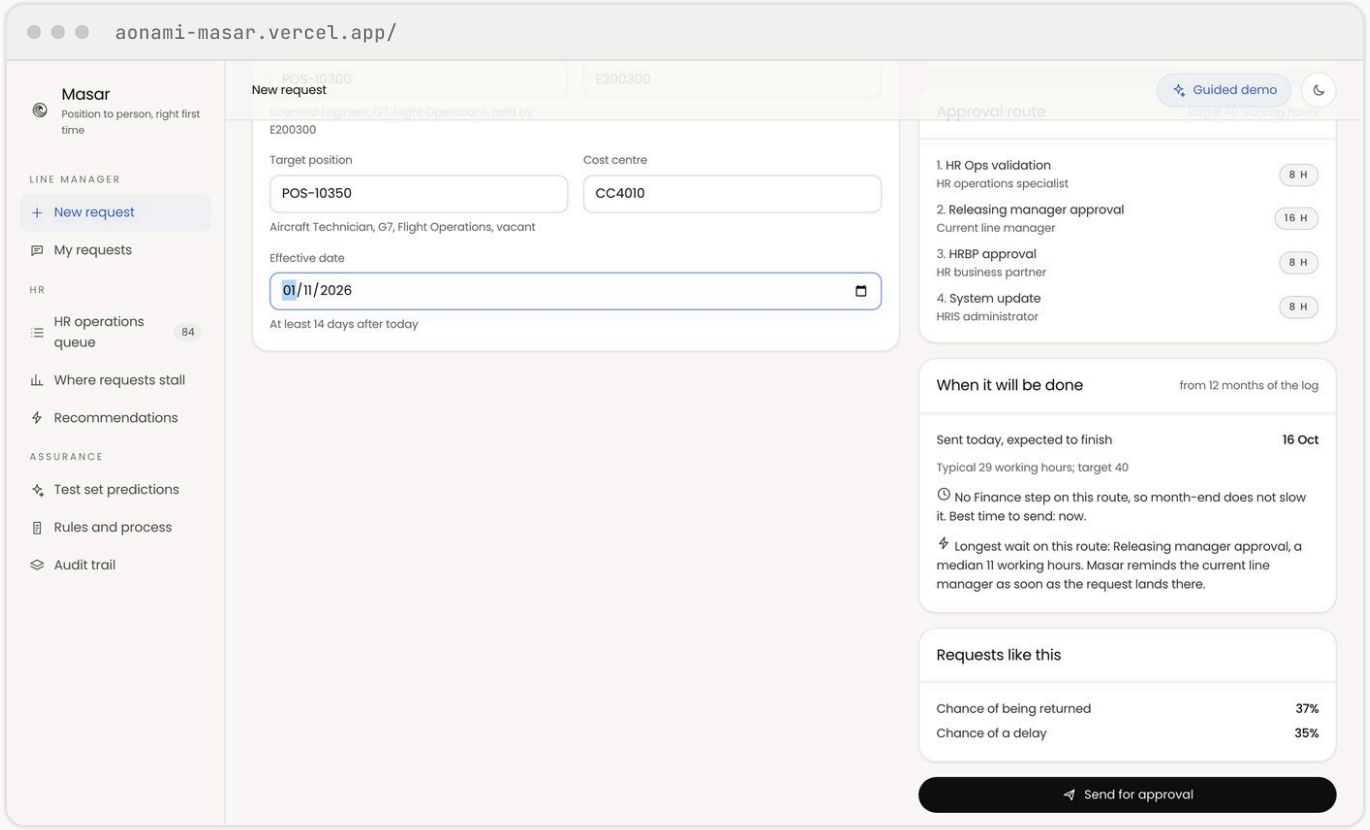
Team variation

HR Ops Airport takes 8.4 hours median to validate against 4.9 for HR Ops Central.

Month-end

Finance steps take 15 hours in the month-end window against 9 otherwise.

Before it is sent, the manager sees when it will be done.



- A **finish date** from the real time each step takes in 12 months of the log, not the target.
- **Best time to send:** if the route has a Finance step, Masar warns about the month-end window, when Finance steps take a median 15 hours against 9 (our sample log).
- The **longest wait** on the route is named, and Masar reminds that approver as soon as the request lands.

WHY IT MATTERS

Managers chase because they cannot see when a request will finish. Showing it before sending, and nudging the slow step, removes the chasing.

Caught at entry, and predicted in flight.

Our own pack in the brief's shape: 3,200 positions, about 5,150 completed requests over 12 months with every step logged, the 19 validation rules and the process map, plus about 150 in-flight requests held out with known outcomes as our own test set.

LIMITS WE DISCLOSE

We wrote this data ourselves. Delays depend more on queues and are harder to predict, so each delay prediction shows its reasons. Rules written in prose need confirming with HR before go-live.

WHAT WE EXPECT ON REAL DATA

Real delays depend on queues, holidays and month-end load we did not fully model. Expect delay accuracy below our 74 to 79% until the predictor is refitted on Emirates' own log in pilot weeks 1 to 2. Return prediction should hold up better, because returns follow the rules.

BENCHMARK ON OUR OWN SYNTHETIC PACK	RESULT
Bad submissions caught at entry	99.1% to 99.8%
Clean submissions wrongly flagged	0
Return prediction, held-out set	96% to 99%, AUC 0.97 to 0.99
Delay prediction, held-out set	74% to 79%, AUC about 0.83
Full analysis of the log, in a browser	under a second

Three generated packs, our own data. The official test set is the real test.

Specific fixes, costed from the log.

R1 CHECK AT ENTRY EFFORT: LOW

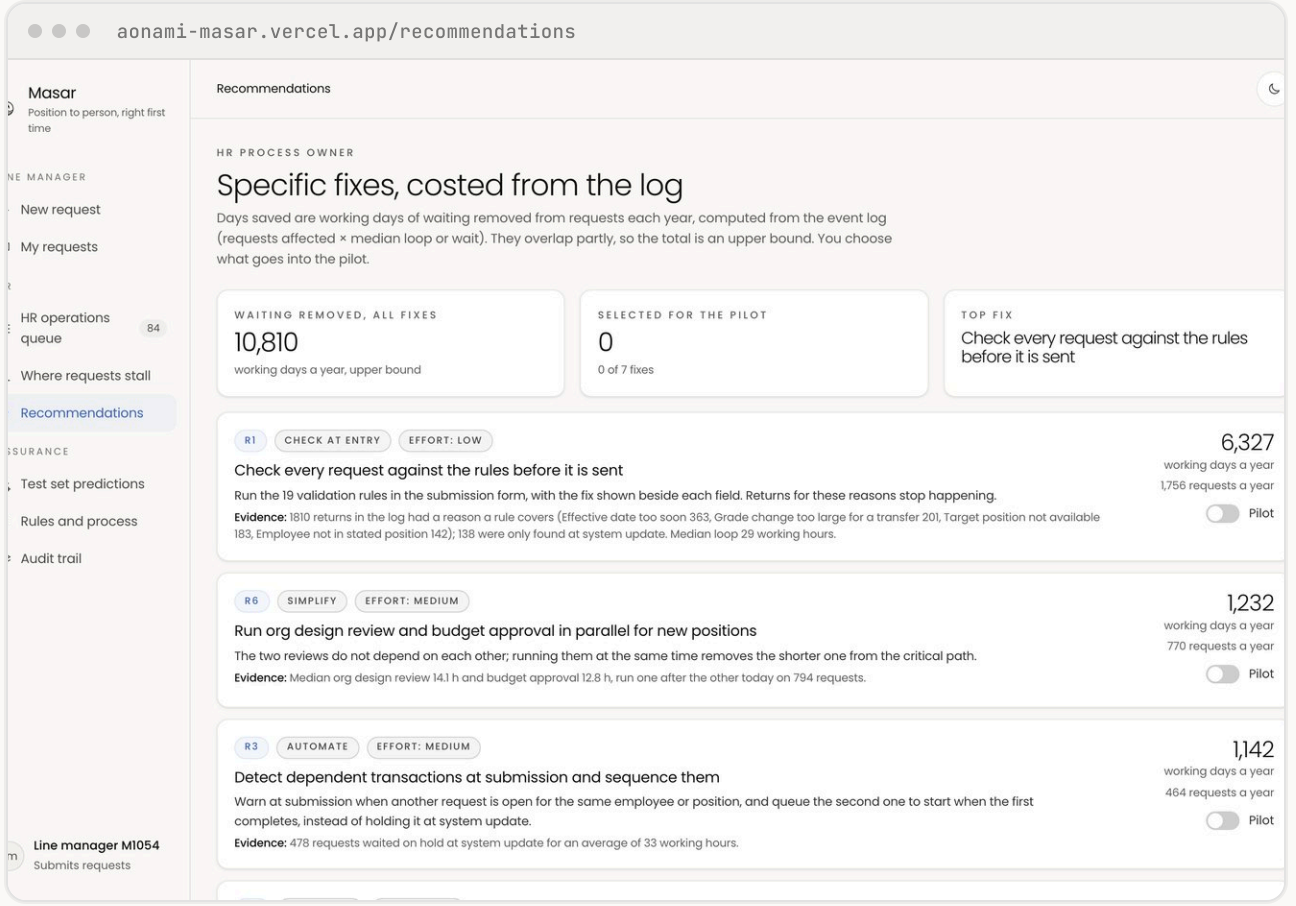
Check every request against the rules before it is sent

About 6,300 working days of waiting removed a year in our sample log, from returns the 19 rules would have caught at submission.

```
days saved = requests affected × median loop or wait, from the log
```

Also: run org design review and budget approval in parallel for new positions (about 1,200 days).

- Fixes overlap, so the total is an upper bound.
- HR chooses which recommendations go into the pilot.



Masar checks, tracks and recommends. People approve.

PEOPLE DECIDE EVERY STEP

The manager decides whether to send; each approval is taken by the named approver.

TRANSPARENT PREDICTOR

A weighting of factors learned from the log; every prediction lists its evidence.

FAIR USE

Predictions are about requests and steps, not the people who submit or approve them.

FULL AUDIT TRAIL

Every submission and decision, with the person and the time.

NO CLOUD AI

Data stays in Emirates' environment.

HR OWNS THE RULES

Rules and the process map are files HR controls.

Designed for the UAE Personal Data Protection Law [2] and the principles of the UAE Charter for the Development and Use of AI [3].

Rules and the process map are files HR owns. Another entity loads its own and everything follows.

1 **HR system**

Read-only feed of P2P transactions, workflow events and the organisation structure.

2 **Masar**

Static web app and command-line runner in a container. No outbound connection.

3 **Single sign-on**

Manager, HR operations, process owner.

4 **Existing workflow**

The checked request is handed to the workflow already in place.

	HOW
Validation rules	Loaded from validation_rules.md; each becomes a live form check.
Process map	Route, approvers and target times from process_map.md.
Event log	Any history in the same shape feeds the insight and the predictor.
Writes back	Never changes a request; hands a checked request to the workflow.

Eight weeks, one directorate, shadow mode first.

WEEKS	WHAT HAPPENS
1 to 2	Load rules, process map and 12 months of the event log; confirm each rule with HR
3 to 4	Shadow mode: check live submissions and predict, without changing the workflow
5 to 8	Live checks for one directorate's managers; status tracking; two fixes piloted
End	Return rate, cycle time and follow-ups before and after; go or no-go

WHAT WE NEED FROM EMIRATES

· The validation rules and the process map

· Twelve months of the P2P event log

· One directorate's managers and its HR operations lead

· A container host inside Emirates with single sign-on

WHAT WE MEASURE

- Return rate before and after
- Cycle time against target
- Manager follow-up emails
- Prediction accuracy on live requests

THE DECISION AT THE END

Go or no-go, on numbers agreed with Emirates before the pilot starts.

Every judging criterion, and where it is answered.

CRITERION	WEIGHT	SLIDES	STRONGEST EVIDENCE
Errors caught and prediction accuracy	30%	7	99.1 to 99.8% of bad submissions caught; 0 clean flagged; return prediction 96 to 99%
Process insight	25%	5, 8	Loops, steps against target, variants, team variation; 7 costed fixes
Manager and HR experience	15%	4, 6	Finish date and best time to send, before the request goes; Form that checks as you type; live status; “where is my request?”
Responsible AI and security	15%	9	People approve every step; transparent predictor; audit trail
Deployability and reuse	15%	10	Rules and the process map are files HR owns

IN ONE LINE

Checked as it is typed, tracked to the next owner, and every fix costed from the log. People approve every step.

TRY IT LIVE

`aonami-masar.vercel.app`

Press Guided demo. Runs on our synthetic data.

Workflow systems, process data and measured prediction.

Jayesh Kumar Rout

JR

CO-FOUNDER, STRATEGY AND PRODUCT

IDfy: KYC and compliance, 60M+ verifications a month

OneCard: onboarding reporting engine for 6 banks

ON THIS BUILD

Owns the pilot plan and the relationship.

Pratik Kumar

PK

CO-FOUNDER, PRODUCT AND GO-TO-MARKET

Stashfin: product manager, credit-report vertical

Aonami: designed the agent approval model in Bridge

ON THIS BUILD

Product flow and the demo; records the video.

Bhavik Munot

BM

FOUNDING ENGINEER, PLATFORM AND DATA

Amazon: compliance-critical logs across 6,000+ AWS accounts

Electronic Arts: real-time pipelines up to 100K TPS

ON THIS BUILD

Data loading, matching and the command-line runner.

Piyush Jha

PJ

FOUNDING ENGINEER, BACKEND AND INFRASTRUCTURE

Ideaflow, Menlo Park: founding engineer, prototype to production

Trell: APIs serving millions of daily users

ON THIS BUILD

Container deployment, single sign-on and connectors.

Harshil Agrawal

HA

ML ENGINEER, EVALUATION AND MLOPS

Nielsen: multi-agent pipeline; 65% less manual analyst time

Monthly retraining and deployment for 50 clients, no manual steps

ON THIS BUILD

Benchmarks: the answer key, precision and recall.

Aonami builds AI agents and decision systems for regulated institutions, with a named person approving every action. Team of 9 across engineering, AI and delivery.

pratik@aonamitech.com

Every number in this deck, and where it comes from.

[1] Challenge brief, Emirates Group on Ignyte Challenges (deadline 16 Oct 2026): judging criteria and weights. Provided to entrants; not public.

[2] UAE Government portal, Data protection laws: Federal Decree-Law No. 45 of 2021 on the Protection of Personal Data. https://u.ae/en/about-the-uae/digital-uae/data/data-protection-laws

[3] CMS Expert Guide, AI regulation in the UAE: UAE Charter for the Development and Use of AI (2024), principles incl. transparency, human oversight, accountability. https://cms.law/en/int/expert-guides/ai-regulation-scanner/united-arab-emirates

[4] Emirates Group, 2025-26 results (Dubai Media Office, 7 May 2026): workforce grew 8% to 130,919; UAE nationals above 4,000. https://www.mediaoffice.ae/en/news/2026/may/07-05/emirates-group-achieves-record-profit-in-2025-26

All people, records and figures on product screens are Aonami's own synthetic data. Emirates Group is named as the challenge host only; no endorsement implied.