Masar: Position to Person, right first time

Checks each P2P request against the rules as the manager fills it in, tracks every request with its next owner, predicts which in-flight requests will be returned or delayed and why, and turns the event log into costed process fixes.

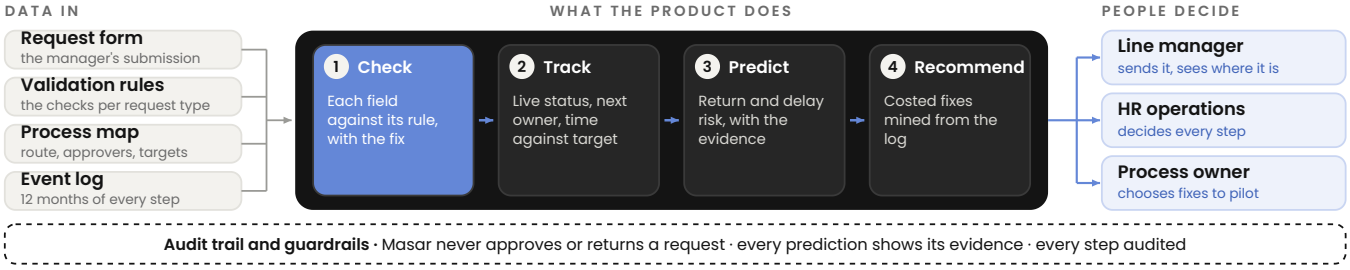
99.5%	95.7%	78.7%	6,293
bad submissions caught at entry, 0 clean flagged	return prediction accuracy, AUC 0.98	delay prediction accuracy, AUC 0.85	working days a year saved by the top fix

Data limitation. The official data pack and test set were not published, so we curated a synthetic sample in the brief's shape (3,200 positions, 5,159 requests, 141 in-flight test requests, seed 2310). Results show the approach on that sample, not on Emirates data. **All findings and files:** github.com/pratikkp24/aonami-masar-emirates-p2p

Try the live prototype
aonami-masar.vercel.app
Hosted, no install. Press **Guided demo**.
Demo video www.loom.com/share/...

THE APPROACH

One loop from submission to insight. The **validation rules** become structured checks that run in the manager's form; the **process map** gives the route, approvers and target times; the **event log** shows where requests stall and loop back. A transparent predictor, learned from the log, scores in-flight requests and lists its evidence. **No cloud AI is used, and Masar never approves or returns a request**: people decide every step.



ERRORS CAUGHT AND PREDICTION ACCURACY

Before the official pack, we built our own in the brief's shape: 3,200 positions, about 5,150 completed requests over 12 months with every step in the event log, the submitted forms (about a third breaking at least one rule), the process map and 19 validation rules, plus about 150 in-flight requests held out with known outcomes as our own test set.

BENCHMARK (THREE GENERATED PACKS)	RESULT
Bad submissions caught at entry	99.1% to 99.8%
Return prediction, in-flight test set	96% to 99% accurate, AUC 0.97 to 0.99
Delay prediction, in-flight test set	74% to 79% accurate, AUC about 0.83

Returns are highly predictable because most come from rule failures; delays depend more on queues and are harder, so each delay prediction shows its reasons (team, month-end, waiting longer than 90% of past requests at the current step, an open dependent transaction).

PROCESS INSIGHT AND RECOMMENDATIONS

From the log, Masar shows each step against its target (median, 90th percentile, share over target), why requests loop back and what each loop costs, the most common paths per request type, and variation between teams. In our sample: 35% of requests are returned at least once; HR Ops Airport takes a median 8.4 working hours to validate against 4.9 for HR Ops Central; Finance steps take 15 hours in the month-end window against 9 otherwise.

TOP RECOMMENDATIONS (SAMPLE LOG)	DAYS A YEAR*
Check every request against the rules before it is sent	about 6,300
Run org design review and budget approval in parallel for new positions	about 1,200
Detect dependent transactions at submission and sequence them	about 1,100
Balance validation work between the two HR Ops teams	about 1,100
Replace free-text clarifications with required fields	about 660

*Working days of waiting removed from requests a year: requests affected × median loop or wait, from the log. Fixes overlap, so the

sum is an upper bound.

RESPONSIBLE AI AND SECURITY STATEMENT

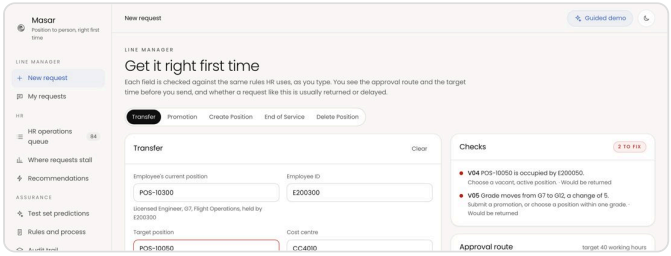
Masar checks, tracks and recommends; people approve. It never approves, returns or changes a request. The manager sees every check and decides whether to send; each approval step is taken by the named approver; HR chooses which recommendations go into a pilot. Every submission and decision is recorded in the audit trail.

No AI model touches the data. The checks are the published rules; the predictor is a transparent weighting of factors learned from the log, and every prediction and recommendation lists its evidence. The synthetic pack was processed only on our own machine and was not uploaded to any service that keeps or trains on data. The hosted demo runs on Aonami's own synthetic sample; files loaded there stay in the viewer's browser.

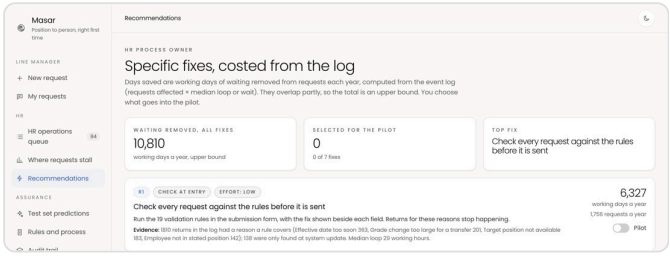
- **Fair use:** predictions are about requests and process steps, not about the people who submit or approve them; team views show where to add support, not individuals to blame.
- **Data protection:** designed for the UAE Personal Data Protection Law (Federal Decree-Law 45 of 2021) and the principles of the UAE Charter for the Development and Use of AI: transparency, human oversight and accountability.
- **Limits we disclose:** results are on our own synthetic benchmark, where we also wrote the data; the official test set is the real test. Rules written in prose need confirming with HR before go-live.

HOW IT RUNS INSIDE EMIRATES' ENVIRONMENT

- **Deployment:** a static web application and a command-line runner, packaged as a container in Emirates' private cloud or on premises. No outbound connection and no cloud AI.
- **Data in:** a read-only feed of P2P transactions and their workflow events from the HR system, plus the organisation structure. The submission form sits in front of the existing workflow and hands a checked request to it.
- **Configuration:** rules and the process map are files HR owns; another entity loads its own and everything follows. Single sign-on roles: manager, HR operations, process owner.



Right first time. Each field is checked against its validation rule as the manager types, with the fix, before the request is sent.



Costed recommendations. Specific fixes from the event log, with requests affected and working days saved a year; HR chooses what to pilot.

PILOT NEXT QUARTER

WEEKS	ACTIVITY
1 to 2	Load rules, process map and 12 months of the event log; confirm each rule with HR
3 to 4	Shadow mode: check live submissions and predict, without changing the workflow; compare with real outcomes
5 to 8	Live checks for one directorate's managers; status tracking; two recommendations piloted
End	Return rate, cycle time and follow-ups before and after; go or no-go

TEAM

Aonami builds AI agents and decision systems for regulated institutions, with a named person approving every action. A team of nine across engineering, AI and delivery; co-founders Pratik Kumar and Jayesh Kumar Rout, with backgrounds in lending, fintech and risk at Stashfin, OneCard and Northern Trust.

Contact: Pratik Kumar · pratik@aonamitech.com · Live demo: aonami-masar.vercel.app